

Bank Marketing Model Performance Analysis

Objective: The objective is to use all this information to predict whether someone will end up saving money with the bank. This helps the bank decide who to focus on when they're trying to get people to save money with them. The Analysis will come across 7 classifiers and each performance for better decision-making. Additionally, the feature importance of various factors will be examined to gain insights into the key drivers influencing the predicted outcomes.

```
# Libraries used
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score
```

✓ 0.0s

Python

Exploratory Data Analysis (EDA)

+ Code

+ Markdown

```
df = pd.read_csv("bank-additional-full.csv", delimiter=';')
print(df.head(5))
```

✓ 0.2s

Python

	age	job	marital	education	default	housing	loan	contact	\
0	56	housemaid	married	basic.4y	no	no	no	telephone	
1	57	services	married	high.school	unknown		no	telephone	
2	37	services	married	high.school	no	yes	no	telephone	
3	40	admin.	married	basic.6y	no	no	no	telephone	
4	56	services	married	high.school	no	no	yes	telephone	

	month	day_of_week	...	campaign	pdays	previous	poutcome	emp.var.rate	\
0	may	mon	...	1	999	0	nonexistent	1.1	
1	may	mon	...	1	999	0	nonexistent	1.1	
2	may	mon	...	1	999	0	nonexistent	1.1	
3	may	mon	...	1	999	0	nonexistent	1.1	
4	may	mon	...	1	999	0	nonexistent	1.1	

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y
0	93.994	-36.4	4.857	5191.0	no
1	93.994	-36.4	4.857	5191.0	no
2	93.994	-36.4	4.857	5191.0	no
3	93.994	-36.4	4.857	5191.0	no
4	93.994	-36.4	4.857	5191.0	no

[5 rows x 21 columns]

Data Description

Column Name	Description
age	Age of the client.
job	Type of job the client has.
marital	Marital status of the client.
education	Level of education of the client.
default	Whether they've had trouble paying debts before. (yes, no, or unknown).
housing	Whether the client has housing loan (yes, no, or unknown).
loan	Whether they have other types of loans. (yes, no, or unknown).
contact	Contact communication type (telephone or cellular).
month	Last contact month of the year. When the bank contacted them.
day_of_week	Last contact day of the week. Which day of the week they were contacted.
duration	Last contact duration, in seconds.
campaign	How many times the bank has contacted them.
pdays	Number of days since the client was last contacted from a previous campaign (-1 means client was not previously contacted).
previous	How many times the bank contacted them before this campaign.
poutcome	Outcome of the previous marketing campaign.
emp.var.rate	How the employment situation is changing. (quarterly indicator).
cons.price.idx	How prices of goods and services are changing.(monthly indicator).
cons.conf.idx	How confident consumers are about the economy.(monthly indicator).
euribor3m	Euribor 3-month rate. A type of interest rate (daily indicator).
nr.employed	How many people are employed. (quarterly indicator).
y	Whether they ended up saving money with the bank or not.(yes or no).

```
df.info()
```

✓ 0.2s

Python

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   41188 non-null  int64
1   job                   41188 non-null  object
2   marital               41188 non-null  object
3   education             41188 non-null  object
4   default               41188 non-null  object
5   housing               41188 non-null  object
6   loan                  41188 non-null  object
7   contact               41188 non-null  object
8   month                 41188 non-null  object
9   day_of_week           41188 non-null  object
10  duration               41188 non-null  int64
11  campaign               41188 non-null  int64
12  pdays                 41188 non-null  int64
13  previous               41188 non-null  int64
14  poutcome              41188 non-null  object
15  emp.var.rate           41188 non-null  float64
16  cons.price.idx         41188 non-null  float64
17  cons.conf.idx          41188 non-null  float64
18  euribor3m              41188 non-null  float64
19  nr.employed            41188 non-null  float64
20  y                      41188 non-null  object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB
```

Univariable Check

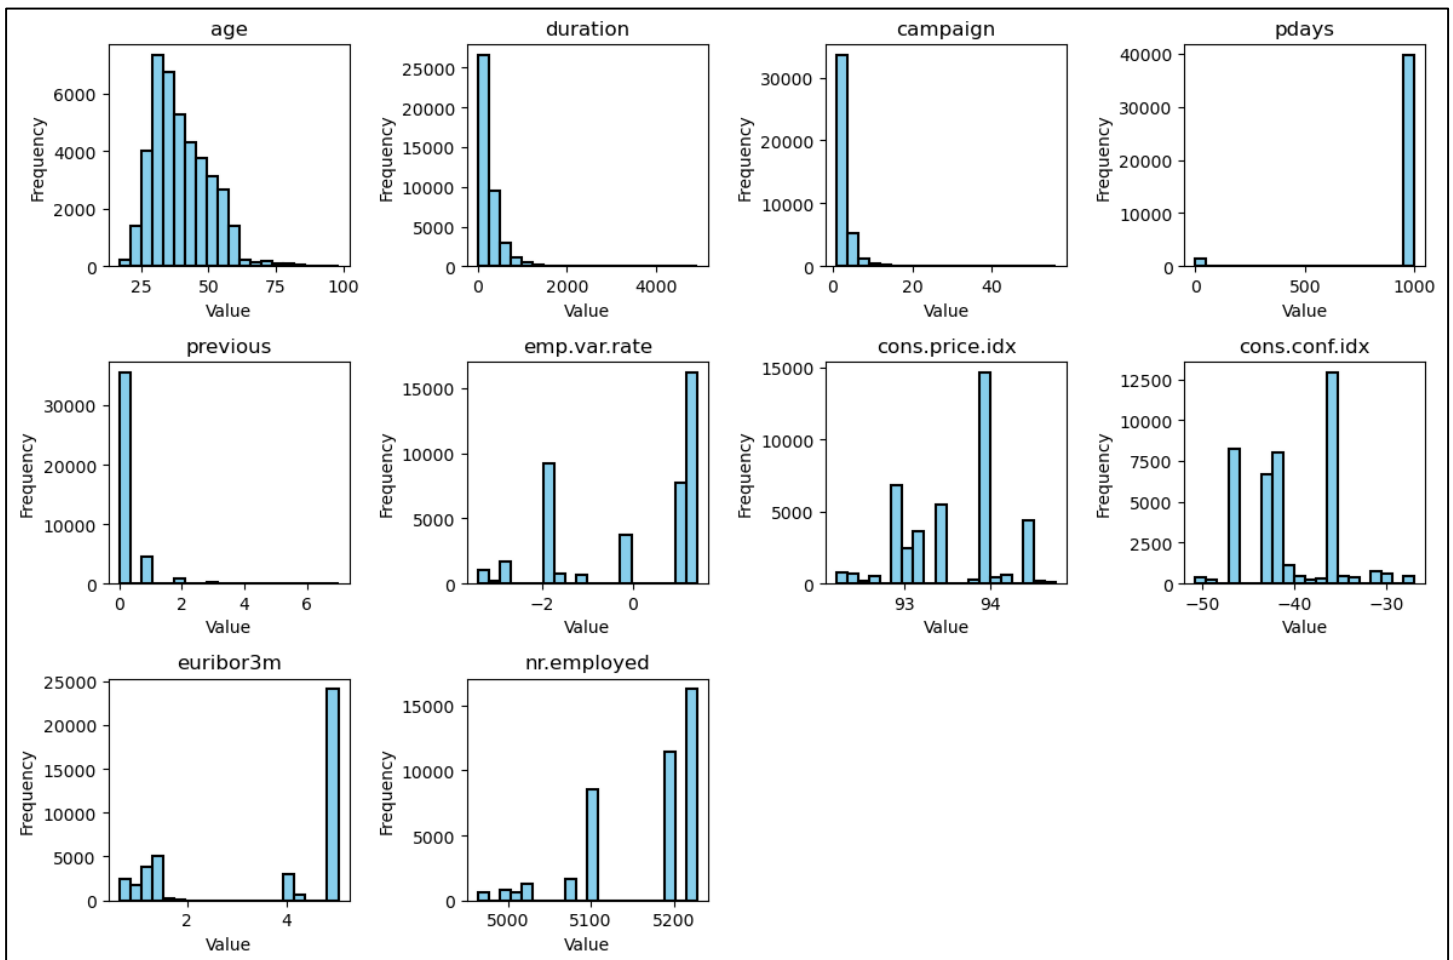
```
# Numerical Columns
numerical_columns = ['age', 'duration', 'campaign', 'pdays', 'previous',
                    'emp.var.rate', 'cons.price.idx', 'cons.conf.idx',
                    'euribor3m', 'nr.employed']

plt.figure(figsize=(12, 8))
for i, column in enumerate(numerical_columns, 1):
    plt.subplot(3, 4, i)
    plt.hist(df[column], bins=20, color='skyblue', edgecolor='black', linewidth=1.5)
    plt.title(column)
    plt.xlabel('Value')
    plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```

✓ 22.9s

Python

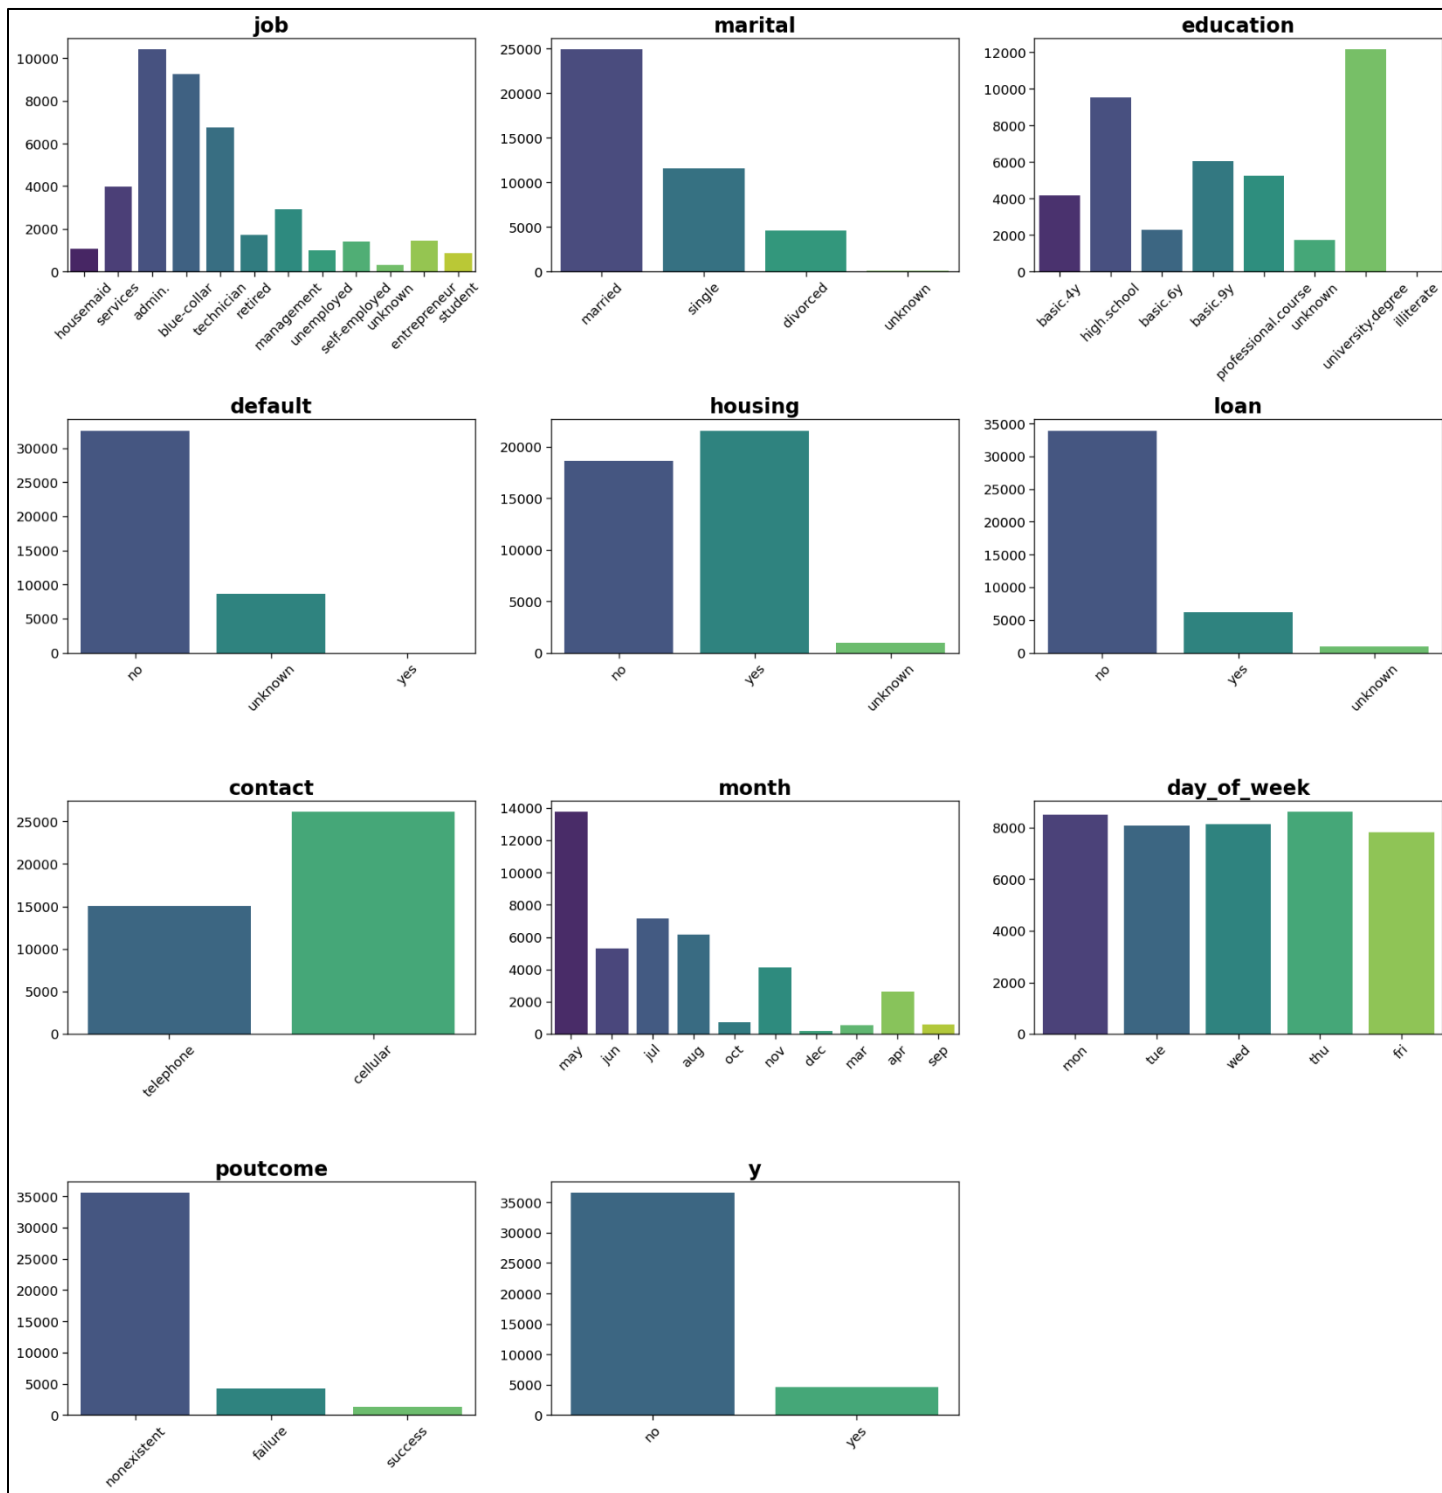


```
# Categorical columns
categorical_columns = ['job', 'marital', 'education', 'default', 'housing', 'loan',
                       'contact', 'month', 'day_of_week', 'outcome', 'y']

plt.figure(figsize=(20, 30))
for i, column in enumerate(categorical_columns, 1):
    plt.subplot(6, 3, i)
    sns.countplot(data=df, x=column, palette='viridis')
    plt.title(column, fontsize=14, fontweight='bold')
    plt.xlabel(None)
    plt.ylabel(None)
    plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

Python



```
summary_stats = df.describe()
summary_stats
```

Python

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	40.02406	258.285010	2.567593	962.475454	0.172963	0.081886	93.575664	-40.502600	3.621291	5167.035911
std	10.42125	259.279249	2.770014	186.910907	0.494901	1.570960	0.578840	4.628198	1.734447	72.251528
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.800000	0.634000	4963.600000
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	93.075000	-42.700000	1.344000	5099.100000
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	93.749000	-41.800000	4.857000	5191.000000
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	93.994000	-36.400000	4.961000	5228.100000
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	94.767000	-26.900000	5.045000	5228.100000

```
# checking missing values
df.isnull().sum()
```

✓ 0.0s

Python

```
age          0
job          0
marital      0
education    0
default      0
housing      0
loan         0
contact      0
month        0
day_of_week  0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
emp.var.rate 0
cons.price.idx 0
cons.conf.idx 0
euribor3m    0
nr.employed  0
y            0
dtype: int64
```

Bivariable Check

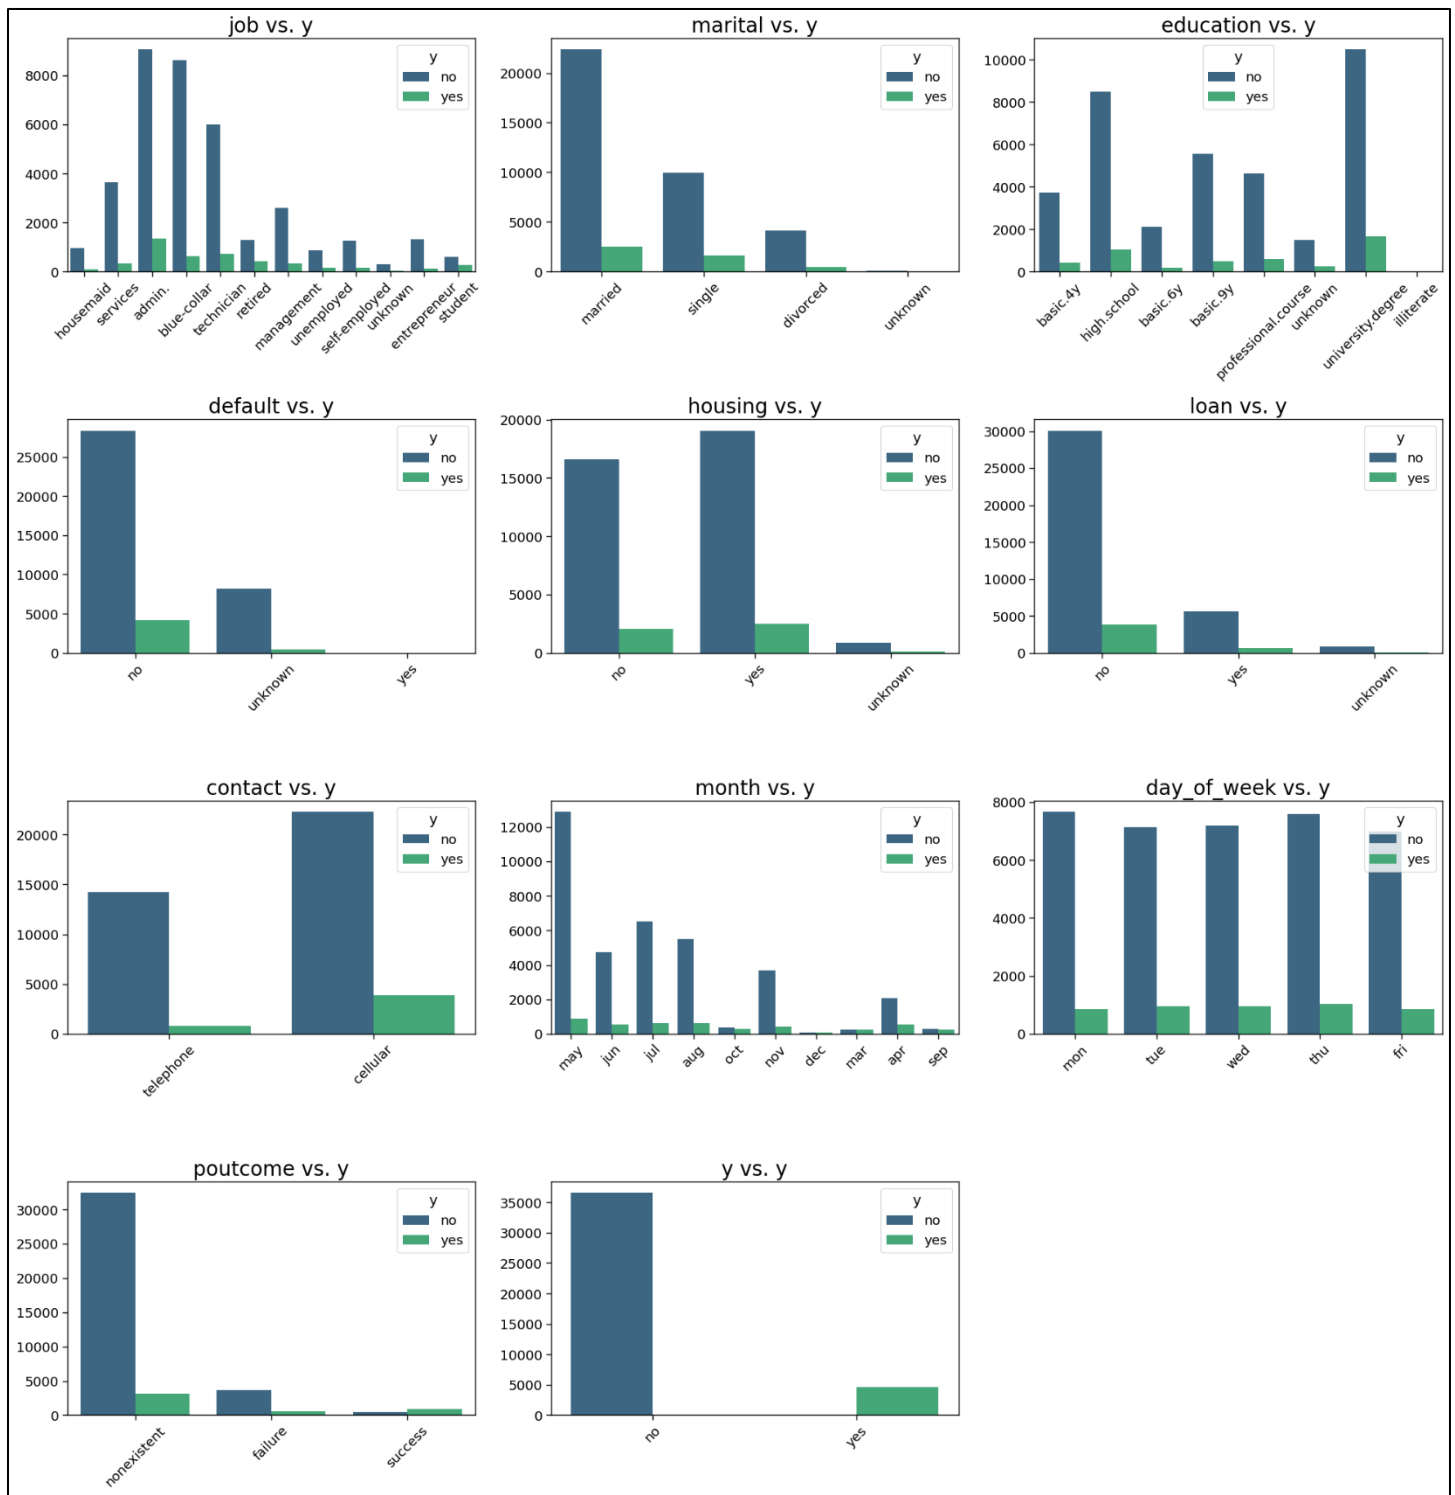
```
# Bivariable check

# Categorical
plt.figure(figsize=(20, 30))
for i, column in enumerate(categorical_columns, 1):
    plt.subplot(6, 3, i)
    sns.countplot(data=df, x=column, hue='y', palette='viridis')
    plt.title(f'{column} vs. y')
    plt.xlabel(None)
    plt.ylabel(None)
    plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

✓ 15.8s

Python

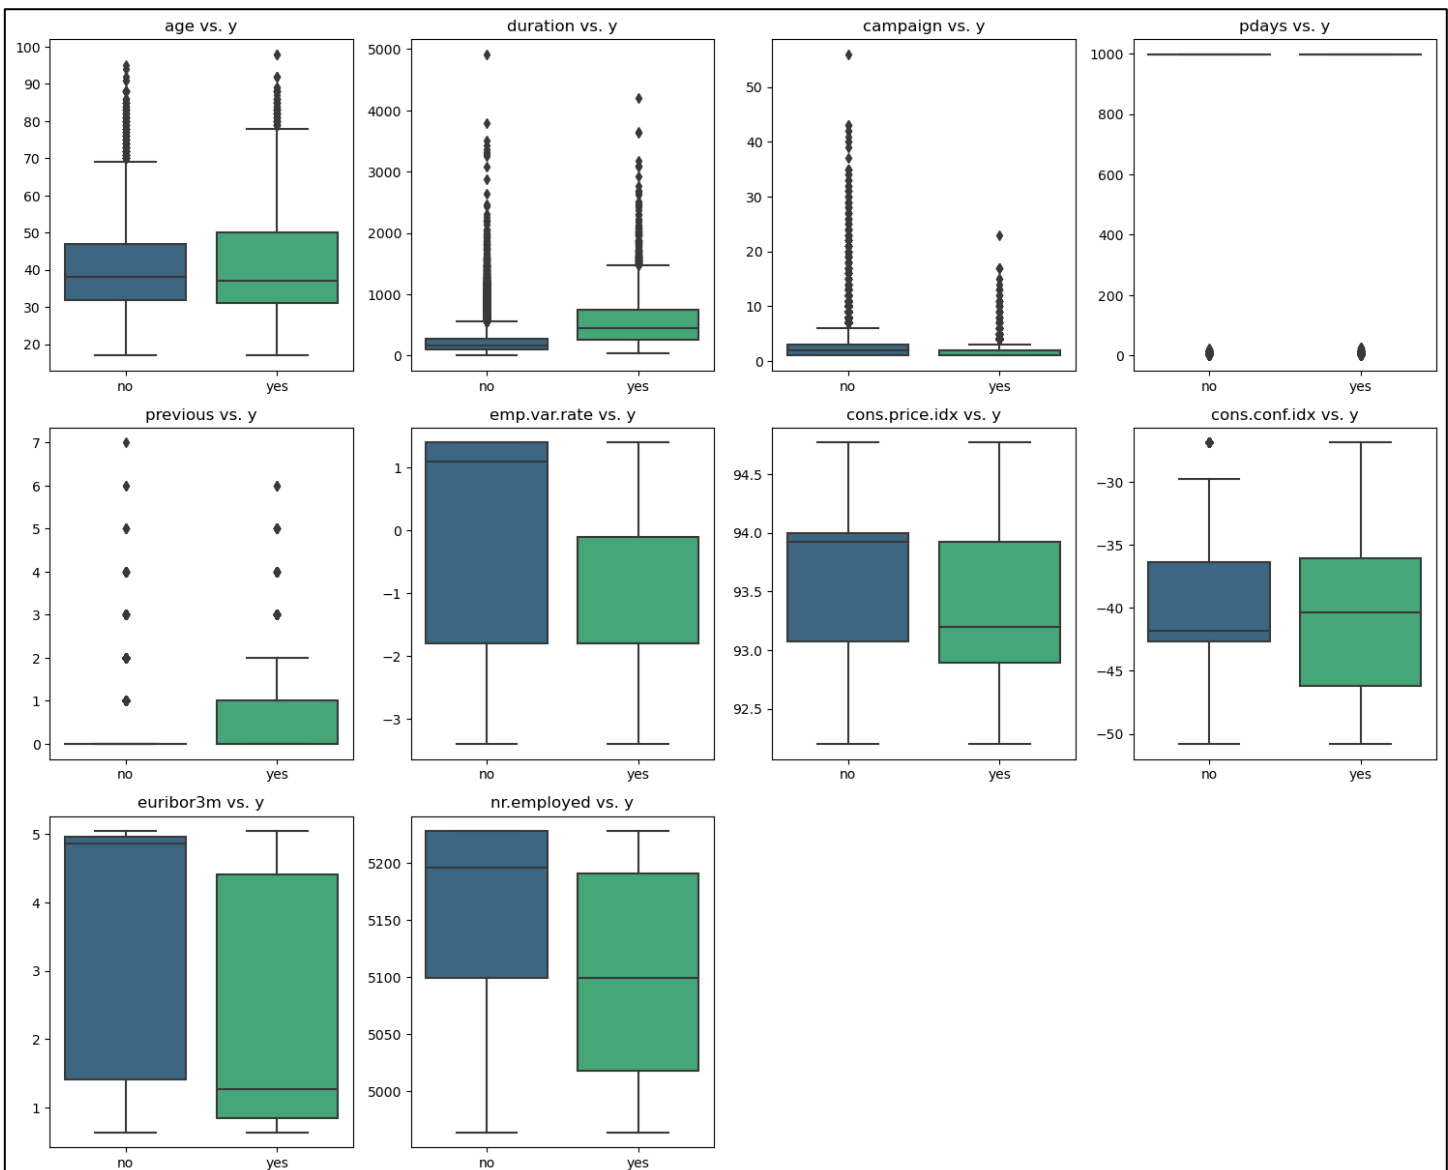


```
# Numerical

plt.figure(figsize=(15, 20))
for i, column in enumerate(numerical_columns, 1):
    plt.subplot(5, 4, i)
    sns.boxplot(data=df, x='y', y=column, palette='viridis')
    plt.title(f'{column} vs. y')
    plt.xlabel(None)
    plt.ylabel(None)

plt.tight_layout()
plt.show()
```

Python




Correlation

```
# correlation

# Compute the correlation matrix
correlation_matrix = df[numerical_columns].corr()

plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", square=True)
plt.title('Correlation Matrix of Numerical Variables')
plt.show()
```

Python



Data Processing

```
# Removing columns with high correlation ( More than 0.95)
variables_to_remove = ['euribor3m', 'nr.employed', 'emp.var.rate']
#drop
df_filtered = df.drop(variables_to_remove, axis=1)

print(df_filtered.head(5))
```

Python

	age	job	marital	education	default	housing	loan	contact	\
0	56	housemaid	married	basic.4y	no	no	no	telephone	
1	57	services	married	high.school	unknown	no	no	telephone	
2	37	services	married	high.school	no	yes	no	telephone	
3	40	admin.	married	basic.6y	no	no	no	telephone	
4	56	services	married	high.school	no	no	yes	telephone	

	month	day_of_week	duration	campaign	pdays	previous	poutcome	\
0	may	mon	261	1	999	0	nonexistent	
1	may	mon	149	1	999	0	nonexistent	
2	may	mon	226	1	999	0	nonexistent	
3	may	mon	151	1	999	0	nonexistent	
4	may	mon	307	1	999	0	nonexistent	

	cons.price.idx	cons.conf.idx	y
0	93.994	-36.4	no
1	93.994	-36.4	no
2	93.994	-36.4	no
3	93.994	-36.4	no
4	93.994	-36.4	no


Python

```
# checking the correlation again

# Numerical Columns
numerical_columns2 = ['age', 'duration', 'campaign', 'pdays', 'previous',
                      'cons.price.idx', 'cons.conf.idx']

# Compute the correlation matrix
correlation_matrix = df_filtered[numerical_columns2].corr()
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", square=True)
plt.title('Correlation Matrix of Numerical Variables')
plt.show()
```

Python



```
# Keeping the column I think is relevant to the model

variables_to_remove2 = ['contact', 'month', 'day_of_week', 'poutcome']
#drop
df_filtered = df_filtered.drop(variables_to_remove2, axis=1)
df_filtered.columns
```

✓ 0.0s

Python

```
Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
      'duration', 'campaign', 'pdays', 'previous', 'cons.price.idx',
      'cons.conf.idx', 'y'],
      dtype='object')
```

Categorical Treatment

+ Code

+ Markdown

-- Marital

```
df_filtered['marital'].value_counts()
```

✓ 0.0s

Python

```
marital
married    24928
single     11568
divorced    4612
unknown      80
Name: count, dtype: int64
```

Marital treatment decision: Remove category "Unknown"

```
df_filtered = df_filtered[df_filtered['marital'] != 'unknown']
df_filtered['marital'].value_counts()
```

✓ 0.0s

Python

```
marital
married    24928
single     11568
divorced    4612
Name: count, dtype: int64
```

-- Housing

[+ Code](#)[+ Markdown](#)

```
df_filtered['housing'].value_counts()
```

✓ 0.0s

Python

```
housing
yes      21541
no       18578
unknown   989
Name: count, dtype: int64
```

Housing treatment decision: Remove category "Unknown"

```
df_filtered = df_filtered[df_filtered['housing'] != 'unknown']
df_filtered['housing'].value_counts()
```

✓ 0.0s

Python

```
housing
yes      21541
no       18578
Name: count, dtype: int64
```

-- Loan

Loan column doesn't need to remove the unknown category because as removing from the others it also removed from the Loan column

-- Education

Create a new category called " basic.education" by replacing the values 'basic.4y', 'basic.6y', and 'basic.9y'

```
df_filtered['education'] = df_filtered['education'].replace(['basic.4y', 'basic.6y', 'basic.9y'],  
                                                             'basic.education')  
df_filtered['education'].value_counts()
```

✓ 0.0s

Python

```
education  
basic.education      12166  
university.degree    11860  
high.school          9281  
professional.course   5112  
unknown              1682  
illiterate            18  
Name: count, dtype: int64
```

-- Age

Age grouping by Equal-wigth Binning : Divide the range of ages into a specified number of equal-width intervals. This approach ensures that each interval has the same width, but it may not capture variations in the distribution of ages.

```
num_bins = 5  
# Create equal-width bins for ages  
df_filtered['age_group'] = pd.cut(df_filtered['age'], bins=num_bins, labels=[f'Group {i+1}'  
                                                                              for i in range(num_bins)])  
  
df_filtered['age_group'].value_counts()
```

✓ 0.0s

Python

```
age_group  
Group 2      19635  
Group 1      12649  
Group 3       7230  
Group 4        508  
Group 5         97  
Name: count, dtype: int64
```

```
# Print the boundaries of each age group
print("Age Group Boundaries:")
print(df_filtered.groupby('age_group')['age'].min())
print(df_filtered.groupby('age_group')['age'].max())
```

Python

```
Age Group Boundaries:
age_group
Group 1    17
Group 2    34
Group 3    50
Group 4    66
Group 5    82
Name: age, dtype: int64
age_group
Group 1    33
Group 2    49
Group 3    65
Group 4    81
Group 5    98
Name: age, dtype: int64
```

```
# removing age from df
df_filtered = df_filtered.drop(columns=['age'])
print(df_filtered.head())
```

✓ 0.0s

Python

```

      job marital      education default housing loan  duration \
0  housemaid married  basic.education      no      no  no      261
1  services married   high.school unknown      no      no      149
2  services married   high.school      no     yes  no      226
3   admin. married  basic.education      no      no  no      151
4  services married   high.school      no      no  yes      307

   campaign  pdays  previous  cons.price.idx  cons.conf.idx   y age_group
0         1     999         0         93.994         -36.4  no  Group 3
1         1     999         0         93.994         -36.4  no  Group 3
2         1     999         0         93.994         -36.4  no  Group 2
3         1     999         0         93.994         -36.4  no  Group 2
4         1     999         0         93.994         -36.4  no  Group 3
```

```
cat_columns = ['job', 'marital', 'education', 'default', 'housing',
               'loan', 'age_group', 'y']

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Encode each categorical column
for column in cat_columns:
    df_filtered[column] = label_encoder.fit_transform(df_filtered[column])

df_filtered.head()
```

Python

	job	marital	education	default	housing	loan	duration	campaign	pdays	previous	cons.price.idx	cons.conf.idx
0	3	1	0	0	0	0	261	1	999	0	93.994	-36.4
1	7	1	1	1	0	0	149	1	999	0	93.994	-36.4
2	7	1	1	0	1	0	226	1	999	0	93.994	-36.4
3	0	1	0	0	0	0	151	1	999	0	93.994	-36.4
4	7	1	1	0	0	1	307	1	999	0	93.994	-36.4

Numerical Treatment

[+ Code](#)[+ Markdown](#)

-- Feature Scaling

```
num_columns = ['duration', 'campaign', 'pdays', 'previous', 'cons.price.idx', 'cons.conf.idx']

# Initialize StandardScaler
scaler = StandardScaler()

df_filtered[num_columns] = scaler.fit_transform(df_filtered[num_columns])
df_filtered.head()
```

Python

	job	marital	education	default	housing	loan	duration	campaign	pdays	previous	cons.price.idx	cons.conf.idx
0	3	1	0	0	0	0	0.010084	-0.566986	0.195436	-0.349162	0.725917	0.88
1	7	1	1	1	0	0	-0.421805	-0.566986	0.195436	-0.349162	0.725917	0.88
2	7	1	1	0	1	0	-0.124881	-0.566986	0.195436	-0.349162	0.725917	0.88
3	0	1	0	0	0	0	-0.414093	-0.566986	0.195436	-0.349162	0.725917	0.88
4	7	1	1	0	0	1	0.187468	-0.566986	0.195436	-0.349162	0.725917	0.88

ML Model

-- Split the data

```
X = df_filtered.drop(columns=['y']) # Features
y = df_filtered['y'] # Target variable

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Print the shapes of the training and testing sets
print("Training set shape:", X_train.shape, y_train.shape)
print("Testing set shape:", X_test.shape, y_test.shape)
```

Python

```
Training set shape: (32095, 13) (32095,)
Testing set shape: (8024, 13) (8024,)
```


- **Logistic Regression**

```
from sklearn.linear_model import LogisticRegression # Library

model = LogisticRegression(random_state=42)

model.fit(X_train, y_train) # Train the model on the training set


y_pred_lr = model.predict(X_test) # Make predictions on the testing set

# Evaluate the model's accuracy
accuracy = accuracy_score(y_test, y_pred_lr)
print("Accuracy:", accuracy)

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred_lr)
plt.figure(figsize=(4, 2))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False, annot_kws={"fontsize":8})
plt.title('Confusion Matrix - LR')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

Python

Accuracy: 0.900049850448654



Logistic Regression achieved an accuracy of 90.0% and demonstrated good performance in classifying the target variable. It correctly classified 6908 instances of the negative class (no) and 314 instances of the positive class (yes).

• Decision Tree

```
from sklearn.tree import DecisionTreeClassifier

dt_classifier = DecisionTreeClassifier()

dt_classifier.fit(X_train, y_train) # Train the decision tree classifier


y_pred_dt = dt_classifier.predict(X_test) # Make predictions on the testing set

accuracy_dt = accuracy_score(y_test, y_pred_dt)
print("Accuracy:", accuracy_dt)

# Calculate confusion matrix
conf_matrix_dt = confusion_matrix(y_test, y_pred_dt)
plt.figure(figsize=(4, 2))
sns.heatmap(conf_matrix_dt, annot=True, fmt='d', cmap='Blues', cbar=False, annot_kws={"fontsize":8})
plt.title('Confusion Matrix - DT')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

Python

Accuracy: 0.8859670987038883



Decision Tree achieved an accuracy of 88.3%. It correctly classified 6604 instances of the negative class and 505 instances of the positive class.

• Random Forest

```
from sklearn.ensemble import RandomForestClassifier # Library

rf_classifier = RandomForestClassifier()
rf_classifier.fit(X_train, y_train) # Train classifier


y_pred_rf = rf_classifier.predict(X_test) # Predict on the test set

# Calculate accuracy
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print("Accuracy:", accuracy_rf)

# Calculate confusion matrix
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(4, 2))
sns.heatmap(conf_matrix_rf, annot=True, fmt='d', cmap='Blues', cbar=False, annot_kws={"fontsize":8})
plt.title('Confusion Matrix - RF')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

Python

Accuracy: 0.9002991026919243



Random Forest achieved an accuracy of 90.2% and demonstrated robust performance. It correctly classified 6807 instances of the negative class and 417 instances of the positive class.

- **Support Vector Machines (SVM)**

```
from sklearn.svm import SVC # Library

svm_classifier = SVC()
svm_classifier.fit(X_train, y_train) # Train classifier


y_pred_svm = svm_classifier.predict(X_test) # Predict on the test set

# Calculate accuracy
accuracy_svm = accuracy_score(y_test, y_pred_svm)
print("Accuracy:", accuracy_svm)

# Calculate confusion matrix
conf_matrix_svm = confusion_matrix(y_test, y_pred_svm)
plt.figure(figsize=(4, 2))
sns.heatmap(conf_matrix_svm, annot=True, fmt='d', cmap='Blues', cbar=False, annot_kws={"fontsize":8})
plt.title('Confusion Matrix - SVM')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

Python

Accuracy: 0.8995513459621136



SVM achieved an accuracy of 90.0%. It correctly classified 6984 instances of the negative class and 234 instances of the positive class.

- **Gradient Boosting**

```
from sklearn.ensemble import GradientBoostingClassifier # Library

gb_classifier = GradientBoostingClassifier()
gb_classifier.fit(X_train, y_train) # Train classifier


y_pred_gb = gb_classifier.predict(X_test) # Predict on the test set

# Calculate accuracy
accuracy_gb = accuracy_score(y_test, y_pred_gb)
print("Accuracy:", accuracy_gb)

# Calculate confusion matrix
conf_matrix_gb = confusion_matrix(y_test, y_pred_gb)
plt.figure(figsize=(4, 2))
sns.heatmap(conf_matrix_gb, annot=True, fmt='d', cmap='Blues', cbar=False, annot_kws={"fontsize":8})
plt.title('Confusion Matrix - NB')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

Python

Accuracy: 0.9086490528414756



Gradient Boosting achieved an accuracy of 90.9% and demonstrated excellent performance. It correctly classified 6839 instances of the negative class and 452 instances of the positive class.

- **K- Nearest Neighbors (KNN)**

```
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_train, y_train) # Train the KNN classifier

y_pred_knn = knn.predict(X_test) # Predict on the testing set


accuracy_knn = accuracy_score(y_test, y_pred_knn)
print("KNN Accuracy:", accuracy_knn)

cm_knn = confusion_matrix(y_test, y_pred_knn)

# Confusion matrix
plt.figure(figsize=(4, 2))
sns.heatmap(cm_knn, annot=True, fmt='d', cmap='Blues', cbar=False, annot_kws={"fontsize":8})
plt.title('Confusion Matrix - KNN')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

Python

KNN Accuracy: 0.892198404785643



KNN achieved an accuracy of 89.2%. It correctly classified 6823 instances of the negative class and 336 instances of the positive class.

- Naïve Bayes

```
from sklearn.naive_bayes import GaussianNB

nb = GaussianNB()

nb.fit(X_train, y_train) # Train the Naive Bayes classifier
y_pred_nb = nb.predict(X_test) # Predict on the testing set


accuracy_nb = accuracy_score(y_test, y_pred_nb)
print("Naive Bayes Accuracy:", accuracy_nb)

cm_nb = confusion_matrix(y_test, y_pred_nb)

# Confusion matrix for Naive Bayes
plt.figure(figsize=(4, 2))
sns.heatmap(cm_nb, annot=True, fmt='d', cmap='Blues', cbar=False, annot_kws={"fontsize":8})
plt.title('Confusion Matrix - Naive Bayes')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```

Python

Naive Bayes Accuracy: 0.889456630109671



Naive Bayes achieved an accuracy of 88.9%. It correctly classified 6748 instances of the negative class and 389 instances of the positive class.

Summarizing this part: Gradient Boosting achieved the highest accuracy among the classifiers tested, followed closely by Random Forest. These models demonstrated robust performance in predicting the target variable, Let's check it further.

Double-checking if the Gradient Boosting and Random Forest are the best among other classifiers comparing through Precision, Recall, and F1 score.

```
# Define the evaluation function
def evaluate_model(y_true, y_pred):
    ... accuracy = accuracy_score(y_true, y_pred)
    ... precision = precision_score(y_true, y_pred, pos_label=1)
    ... recall = recall_score(y_true, y_pred, pos_label=1)
    ... f1 = f1_score(y_true, y_pred, pos_label=1)
    ... return accuracy, precision, recall, f1

# Evaluate each classifier
classifiers = {
    ... "Logistic Regression": y_pred_lr,
    ... "Decision Tree": y_pred_dt,
    ... "Gradient Boosting": y_pred_gb,
    ... "Random Forest": y_pred_rf,
    ... "Support Vector Machine (SVM)": y_pred_svm,
    ... "K-Nearest Neighbors (KNN)": y_pred_knn,
    ... "Naive Bayes": y_pred_nb
}

for clf_name, y_pred in classifiers.items():
    ... accuracy, precision, recall, f1 = evaluate_model(y_test, y_pred)
    ... print(f"{clf_name}:")
    ... print(f"Accuracy: {accuracy:.4f}")
    ... print(f"Precision: {precision:.4f}")
    ... print(f"Recall: {recall:.4f}")
    ... print(f"F1 Score: {f1:.4f}")
```

Python

Logistic Regression:

Accuracy: 0.9000
Precision: 0.6267
Recall: 0.3380
F1 Score: 0.4392

Decision Tree:

Accuracy: 0.8860
Precision: 0.5070
Recall: 0.5436
F1 Score: 0.5247

Gradient Boosting:

Accuracy: 0.9086
Precision: 0.6384
Recall: 0.4865
F1 Score: 0.5522

Random Forest:

Accuracy: 0.9003
Precision: 0.5915
Recall: 0.4489
F1 Score: 0.5104

Support Vector Machine (SVM):

Accuracy: 0.8996
Precision: 0.6783
Recall: 0.2519
F1 Score: 0.3673

K-Nearest Neighbors (KNN):

Accuracy: 0.8922
Precision: 0.5526
Recall: 0.3617
F1 Score: 0.4372

Naive Bayes:

Accuracy: 0.8895
Precision: 0.5285
Recall: 0.4187
F1 Score: 0.4673

Feature Importance

Feature importance analysis is a technique used to determine the relative importance of each feature in predicting the target variable.

Extracting feature importance from Gradient Boosting and Random Forest classifiers, considering they exhibited the best performance

-- Gradient Boosting Feature Importance

```
feature_importances_gb = gb_classifier.feature_importances_  
  
# Display feature importances  
feature_importance_df_gb = pd.DataFrame  
({'Feature': X_train.columns, 'Importance': feature_importances_gb})  
feature_importance_df_gb = feature_importance_df_gb.sort_values(by='Importance', ascending=False)  
print("Gradient Boosting Feature Importance:")  
print(feature_importance_df_gb)
```

Python

Gradient Boosting Feature Importance:

	Feature	Importance
6	duration	0.449408
8	pdays	0.211973
11	cons.conf.idx	0.192609
10	cons.price.idx	0.114420
12	age_group	0.009168
9	previous	0.008324
3	default	0.005169
2	education	0.003643
7	campaign	0.003206
0	job	0.001631
1	marital	0.000203
5	loan	0.000139
4	housing	0.000106

-- Random Forest Feature Importance

```
feature_importances_rf = rf_classifier.feature_importances_  
  
# Display feature importances  
feature_importance_df_rf = pd.DataFrame  
({'Feature': X_train.columns, 'Importance': feature_importances_rf})  
feature_importance_df_rf = feature_importance_df_rf.sort_values(by='Importance', ascending=False)  
print("Random Forest Feature Importance:")  
print(feature_importance_df_rf)
```

✓ 0.8s

Python

Random Forest Feature Importance:

	Feature	Importance
6	duration	0.383752
11	cons.conf.idx	0.107900
10	cons.price.idx	0.085974
8	pdays	0.080620
0	job	0.068586
7	campaign	0.060461
2	education	0.049182
12	age_group	0.045446
9	previous	0.033405
1	marital	0.031473
4	housing	0.025164
5	loan	0.017168
3	default	0.010871

- **Gradient Boosting**

Duration: This feature has the highest importance, indicating that the duration of the call has a significant impact on the outcome.

Pdays: The number of days that passed after the client was last contacted from a previous campaign is also a crucial factor.

Cons.conf.idx and Cons.price.idx: These are economic indicators, suggesting that the overall economic context plays a role.

Age Group and Previous Contacts: These features have relatively lower importance but still contribute to the model.

- **Random Forest**

Duration: Similarly, the duration of the call is the most critical predictor in the Random Forest model.

Cons.conf.idx and Cons.price.idx: Economic indicators remain significant in this model as well.

Job and Campaign: Job type and number of contacts during this campaign also have notable importance.

Education and Age Group: These features also contribute significantly to the model's predictions.

- **Conclusion to this part**

Both models highlight the importance of the call duration and economic indicators (cons.conf.idx and cons.price.idx). Other factors such as job type, education level, and age group also play essential roles in predicting the outcome of the marketing campaign. Overall, these insights can guide marketing strategies to focus on specific customer demographics and tailor communication strategies based on economic conditions and call duration.

ROC Curve form Gradient Boosting and Random Forest

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt


# Compute predicted probabilities for Gradient Boosting and Random Forest
y_prob_gb = gb_classifier.predict_proba(X_test)[:, 1]
y_prob_rf = rf_classifier.predict_proba(X_test)[:, 1]

# Compute ROC curve and ROC area for Gradient Boosting
fpr_gb, tpr_gb, _ = roc_curve(y_test, y_prob_gb)
roc_auc_gb = auc(fpr_gb, tpr_gb)

# Compute ROC curve and ROC area for Random Forest
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_prob_rf)
roc_auc_rf = auc(fpr_rf, tpr_rf)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_gb, tpr_gb, color='blue', lw=2, label=f'Gradient Boosting (AUC = {roc_auc_gb:.2f})')
plt.plot(fpr_rf, tpr_rf, color='red', lw=2, label=f'Random Forest (AUC = {roc_auc_rf:.2f})')
plt.plot([0, 1], [0, 1], color='gray', lw=1, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```

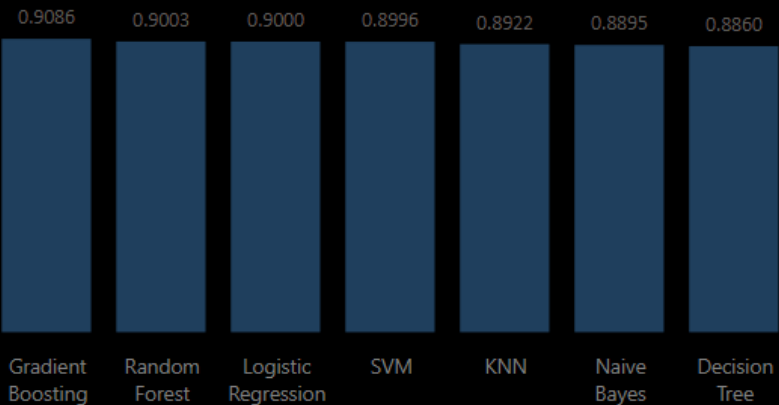
Python






BANK MARKETING MODEL PERFORMANCE DASHBOARD

Accuracy Comparison



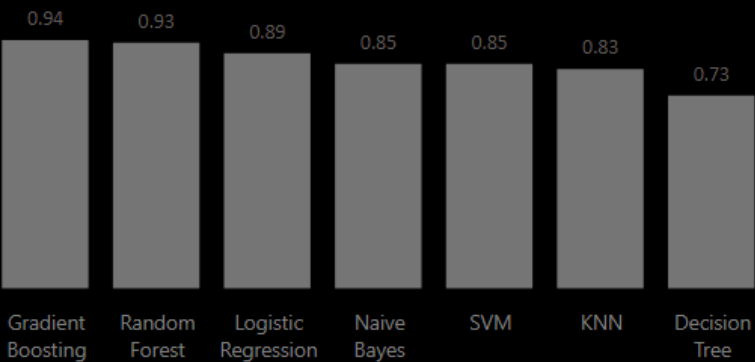
Performance Metrics Overview



Precision-Recall Tradeoff



AUC Comparison



Conclusion:

After integrating all classifiers into the Power BI dashboard and comparing their performance, it's evident that Gradient Boosting and Random Forest consistently outperform other classifiers. This conclusion is drawn from various performance metrics such as Accuracy, Precision, Recall, F1 Score, and AUC. Moreover, analyzing the feature importance reveals that the duration of the call holds significant importance for both classifiers. This insight suggests that investing time in client calls could be an effective strategy for the bank.